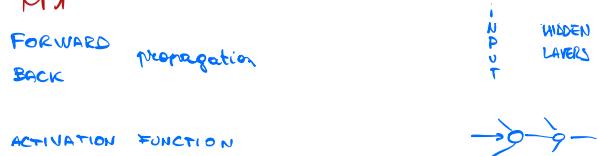


KERAS

M1

FORWARD PROPAGATION
BACK

ACTIVATION FUNCTION



$$\text{SIGMOID} \Rightarrow a = f(x) = \frac{1}{1 + e^{-x}}$$

M2

GRADIENT DESCENT

- ↳ find minimum
- ↳ select step learning rate

BACK PROPAGATION

- ↳ weight & biases OPTIMIZATION
- PROPAGATE ERROR BACK
- ERROR TO UPDATE W

VANISHING

- no no ACTIVATION
- fast at first
- slow at end

ACTIVATION

1. BINARY
 2. LINEAR
 3. SIGMOID
 4. Hyperbolic Tangent → → Vanishing to 0 to 1
 5. ReLU
 6. Leaky ReLU
 7. Softmax
- No vanishing
- Vanishing on deep NN
- probabilities []

COST FUNCTION

$$S = \sum_{i=1}^m (z_i - w x_i - b)^2$$

INITIALIZE weight gradient second
upgrading $w = w - \eta \cdot \frac{\partial S}{\partial w}$

M3

LIBS

- ↳ TF (PRODUCTION) Google
- ↳ Keras HIGH LEVEL API on top of TF by GOOGLE → QUICK
- ↳ PyTorch GPU trained by TF ++ CUSTOMIZATION → DEEP

KERAS

EX1 CONCRETE

```
import keras
from k import Sequential
from k import Dense
```

model = Sequential()
m_cols = concrete_data.shape[1]

L1 model.add(Dense (neurons, activation, input_shape ()))

Lm

model.compile(optimizer, loss) mean squared REGRESSION

EX-2 CAR DATA NOT encoded

```

import keras as K
from K.models import Sequential
from K.layers import Dense
from K.utils import to_categorical

model = Sequential()
n_cols = car_data.shape[1]
target = to_categorical(target)

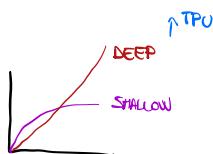
model.add(Dense(neurons, activation, input_shape()))
model.add(Dense(5, "relu"))
model.add(Dense(1, "softmax")) → 0 to 1 % output

model.compile(optimizer, loss, metrics [accuracy])
model.fit(predictors, target, epochs)
model.predict(test_data)
    1 in 4
  
```

M4 Deep Learning

Shallow NN → INPUT VECTOR

Deep NN → large amounts of data



Convolutional NN (CNN)

→ IMAGES AS INPUTS

IMAGE RECOGNITION, CV



CONV LAYERS (w/o RELU)

L. POOLING make foreach section
only

model = Sequential

input_shape = (128, 128, 3)

model.add(Conv2D(16, kernel_size, strides, padding, activation))

◦ MaxPooling2D

◦ Conv2D (32)

◦ MaxPooling2D

◦ Flatten

◦ Dense (100)

◦ Dense



Recurrent NN (RNN)

input output before new input

TEXT, GENOME, FINANCE

IMAGE GENERATION

For short term model

11	12
2	13
3	14
4	15
5	16
6	17
7	18
8	19
9	20
10	

AUTO ENCODERS

data denoise

can't be reused

uses Backprop.

auto minimize

1. fix imbalanced dataset

2. estimating missing values

3. automatic feature extraction

$$a_{11} = e \\ z_{11} = 0.51 \\ a_{21}$$

$$z_{22} = \\ a_{12} = e \\ a_{22}$$

$$\frac{\partial E}{\partial s} = \frac{\partial E}{\partial a_{21}}$$

$$(-f - a_{22}) \cdot (a_{22}(1-a_{22})) \cdot (a_{12})$$

PYTORCH BASICS

ML

TENSORS

DATA → TENSOR → NN inputs
 vectors
 array
 matrix
 w/ NumPy

torch.	float 32	.float
	float 64	.double
	float 16	.half
	int 8	.int8
	uint8	.uint8
	int 16	.int16
	int 32	.int32
	int 64	.int64
	bool	.bool
		.short
		.int
		.long

1D TENSORS

```
import torch
a = torch.tensor([0, 1, 2, 3, 4])
a.dtype      → type of data
a.type       → type of tensor
```

SPECIFIC TYPE `tensor.FloatTensor`
 CHANGE TYPE `a = a.type(torch.FloatTensor)`

`a.size`
`a.ndimension`

convert 1D to 2D

`[0...6]`
 to $\begin{bmatrix} 0 \\ \vdots \\ 6 \end{bmatrix}$ `a.view(6, 1)`

^{tensor}
 PYTORCH \Rightarrow NUMPY
`mp_x = np.array([0 ... 5])`
`torch_x = torch.from_numpy(mp_x)`
`back_to_mp = torch_x.numpy()`

INDEXING
`torch_x = torch.tensor([0 .. 5])`
`torch_x[0] = 2 → [2 .. 5]`

Slicing
 Extract `elt = torch_x[3:5]`
 To change `torch_x[3:5] = torch.tensor([6, 8, 10])`

OPERATIONS

VECTOR ADDITION (same type) `torch_1 + torch_2`
 VECTOR × SCALAR `3 * torch`
 PRODUCT 2 VECTORS `torch_1 * torch_2`
 (dot product) `torch_1.dot(torch_2)`
 ADD CONSTANT `torch_1 + #`

APPLY FUNCTIONS `torch_1.mean()`
 .max()
 UNIVERSAL FUNCTION
 .sin
 .cos

`torch.linspace(start, end, steps)`

2D TENSORS

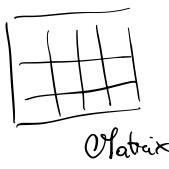


Table & Images
 or Grid



MATRIX
 MULT

`data = [[r1], [r2], [r3]]`
`tensor = torch.tensor(data)`
 .shape()
 .size()
 .numel()

VECTOR
 MATRIX MULT → mm()

DERIVATIVE as graph

$$f = x^2$$

↳ backward derivative of y

data	2
grad_fn	None
grad	8
is_leaf	True
requires_grad	True

$$\frac{d}{dx} (x^2) = 2x$$

$$\begin{aligned} z(x) &= x^2 + 2x + 1 \\ z(2) &= 2^2 + 2(2) + 1 = 9 \Rightarrow \frac{dz}{dx} = 2x + 2 \\ \frac{d^2z}{dx^2} &= 2(2) + 2 = 6 \end{aligned}$$

$$f(u, v) = 4v + u^2$$

$$u = \text{torch.tensor}(1, \text{requires_grad=True})$$

$$v = \text{torch.tensor}(2, \text{requires_grad=True})$$

$$f = u * v + u * 2$$

$$f.backward()$$

$$u.grad$$

$$v.grad$$

$$f(u, v) = 4v + u^2$$

$$\frac{\partial f(u, v)}{\partial u} = v + 2u \Rightarrow$$

$$\frac{\partial f(u, v)}{\partial v} = u$$

$$\frac{d}{dx} (x^2) = 2x$$

data	4
grad_fn	None
grad	8
is_leaf	False
requires_grad	True

SIMPLE DATASET

```
class Element_DS(Dataset):
    def __init__(self, length, transform):
        self.length = length
        self.transform = transform

    def __getitem__(self, index):
        return self
```

METHODS REQUIRED for COMPATIBILITY

- init -
- get item -
- len -

APPLY

multiple data Transform
using Compose method
transforms.Compose [f₁(), f₂(), ..., f_n()]

DATASETS

```
from torch.utils.data import Dataset, DataLoader
```

ZALANDO FASHION
+ category image
0 "t-shirt" 1
1 "trousers" 2
2 "pullover" 3
3 "dress" 4
4 "coat" 5
5 "sandal" 6
6 "shirt" 7
7 "sneaker" 8
8 "bag" 9
9 "ankle boot" 10

```
class Dataset(Dataset):
    def __init__(self, csv_file, data_dir, transform=None):
        self.csv_file = csv_file
        self.data_dir = data_dir
        self.transform = transform
        self.labels = []
        self.images = []

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, idx):
        img_name = os.path.join(self.data_dir, self.csv_file.iloc[idx, 0])
        image = Image.open(img_name)
        label = self.csv_file.iloc[idx, 1]
        if self.transform:
            image = self.transform(image)
        return image, label
```

```
import torchvision.transforms as transforms
transforms.CenterCrop(20)
transforms.ToTensor()
Compose ->
= transforms.Compose([
    transforms.CenterCrop(20),
    transforms.ToTensor()])
dataset = Dataset(csv_file='fashion-mnist_train.csv',
                  data_dir='fashion-mnist_train',
                  transform=5)
```

USE KNOWN DATASETS

```
import torchvision.datasets as dsds
```

```
dataset = dsds.MNIST('root = './data', train=False, download=True, transform=transforms.ToTensor())
```

convert image to tensor

M2

LINEAR REGRESSION 1D

feature x
target y

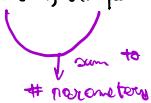
$$y = b + w \cdot x$$

↓
BAS
↑
WEIGHT

DATA → TRAIN
|
FORWARD

CLASS LINEAR

```
from torch.nn import Linear
model = Linear(in_features=1, out_features=1)
y = model(x)
```



TRAIN

$$\mathcal{D} = \{(x_1, y_1), \dots, (x_m, y_m)\}$$

always error \rightarrow Gaussian \sim deviation from the line
NOISE ASSUMPTION

↳ MINIMIZZO la FUNZIONE
mean squared error

LOSS

distance from module estimate

Loss

$$y_i \rightarrow \text{loss}(y_i) \leftarrow \hat{y}_i = w \cdot x_i$$

↓
#

0 → GOOD
HIGH → BAD

MINIMIZE THE ERROR \rightarrow slope

provides a number that helps estimate how good the prediction is.

GRADIENT DESCENT (OR BATCH G.D.)

↳ method find min of a function

considering loss function

Learning rate is the step

if: too big we can miss the minimum

if: too low \rightarrow small

WE CAN STOP TUNING GRADIENT DESCENT

initial value is random \rightarrow then tune

import torch.nn as nn
class LR(nn.Module)
def __init__(self, in_size, out_size):
super(LR, self).__init__()
self.linear = nn.Linear(in_size, out_size)

CUSTOM MODULE
LINEAR

def forward(self, x):
out = self.linear(x)
return out

model = LR(1, 1)
no need to load all data

model.state_dict()
python dict keys & values

$$\text{loss}(w) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$\hat{y}_i = w \cdot x_i$
estimate

? how to find best w ?
↳ minimum

COST FUNCTION

↳ can be iterated

$$\frac{\partial \text{loss}}{\partial w} = -\frac{2}{N} \sum_{n=1}^N (y_n - w \cdot x_n) x_n$$

check if learning rate is + or - during tuning

All samples are called Batch

LINEAR REGRESSION Pytorch

x is torch \rightarrow numpy \rightarrow visualization w/ matplotlib

$$w = \text{torch.tensor}(-10.0, \text{ requires_grad=True})$$

$$x = \text{torch.arange}(\text{start}, \text{stop}, \text{range}).\text{view}(-1, 1)$$

$$f = -3 * x$$

$$\text{forward} \rightarrow \text{return } w * x \approx \hat{y} = w * x$$

$$\text{criterion} = \text{torch.mean}((\hat{y} - y)^2) \approx \frac{1}{N} \sum_{n=1}^N (\hat{y}_n - y_n)^2$$

going through data = EPOCH

for epoch

\hat{y} = forward(x)

loss = criterion(\hat{y} , y)

loss.backward()

$w.\text{data} = w.\text{data} - \eta \text{ step } w.\text{grad}.data$

$w.\text{grad}.data.\text{zero_}()$

cost.append(loss.item()) # displayable

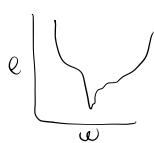
TRAIN SLOPE & BIAS

$$l(w, b) = \frac{1}{N} \sum_{n=1}^N (y_n - w x_n - b)^2$$

pytorch



find minimum of the surface



GRADIENT points in the direction of the next iteration

epoch

\hat{y} = forward

loss = criterion(\hat{y} , y)

loss.backward()

$w.\text{data} = w.\text{data} - \eta \times w.\text{grad}.data$

$w.\text{grad}.data.\text{zero_}()$

$b.\text{data} = b.\text{data} - \eta \times b.\text{grad}.data$

$b.\text{grad}.data.\text{zero_}()$

OPTIMIZATION \rightarrow Stochastic Gradient Descent

minimize 1 approximation at time

η = #

LOSS = []

COST = []

for epoch in range

total = 0

for x, y in zip(x, y)

\hat{y} = forward(x)

loss = criterion(\hat{y} , y)

loss.backward()

$w.\text{data} = w.\text{data} - \eta \times w.\text{grad}.data$

$b.\text{data} = b.\text{data} - \eta \times b.\text{grad}.data$

$w.\text{grad}.data.\text{zero_}()$

$b.\text{grad}.data.\text{zero_}()$

loss.append(loss.item())

total += loss.item()

cost.append(total)

OPTIMIZATION → mini-Batch Descent

process larger data, we use few samples at time

$$? = \# \text{ iterations} = \frac{\text{Training size}}{\text{Batch size}}$$

OPTIMIZATION in PyTorch

DATASET

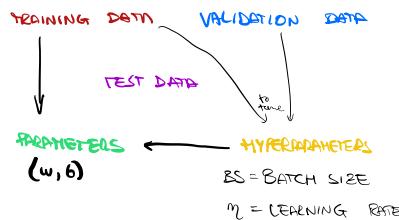
CUSTOM MODULE

COST MODULE ⇒

```
model = LR(1, 1)
from torch import nn, optim
optimizer = optim.SGD(model.parameters(), lr=η)
```

TRAINING, VALIDATION & TEST SPLIT

```
for epoch in range(100)
    for x, y in trainloader
        yhat = model(x)
        loss = criterion(yhat, y)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        w.data = w.data - η * w.grad.data
        b.data = b.data - η * b.grad.data
```



we test multi models
 ↳ Then decide where cost
 is less and choose that model

SAVING MODEL USING CHECKPOINTS

```
class Dataset
    train_data = Data()
    val_data = Data(train=False)           → IBM_013
```

EARLY STOPPING

We save a model, based on cost results

```
dataset
model
TRAINING
loss_val = criterion(model(val_data.x), val_data.y).item()
if loss_val < min_loss:
    value = epoch
    min_loss = loss_val
torch.save(model.state_dict(), 'best-model.pt')
# We save if value is less
```

MODEL + CHECKPOINT

```
checkpoint["epoch"]
    ["model-state-dict"] = model.state_dict()
    ["optimizer-state-dict"] = optimizer.state_dict()
    ["loss"] = loss
torch.save(checkpoint, checkpoint_path)
```

checkpoint = torch.load(checkpoint_path)

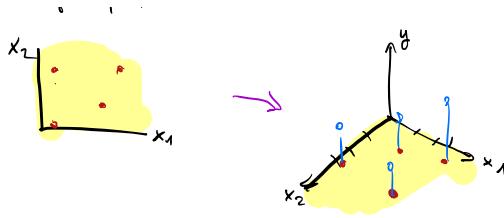
M4

MULTIPLE LINEAR REGRESSION (MLR)

$$\hat{y}_j = b_0 + w_1 x_1 + w_2 x_2 + w_3 x_3 + w_4 x_4$$

$$b_0 = \text{bias}$$

$$w_1 = \text{weight of } x_1$$



$x \cdot w \rightarrow$ DOT PRODUCT

2D TENSORS / MATRIX

our CLASS LINEAR

```
from torch.nn import Linear  
torch.manual_seed(1)  
model = Linear(in_features=2, out_features=1)  
list(model.parameters())
```

MULTIPLE LINEAR REGRESSION

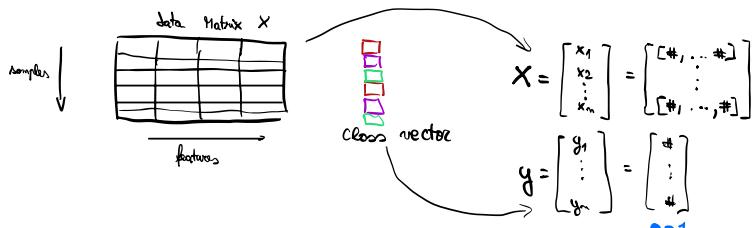
$$\hat{y}_1 = b_{01} + w_{11}x_1 + w_{12}x_2 + w_{13}x_3 + w_{14}x_4$$

$$\begin{bmatrix} \hat{q}_1 & \hat{q}_2 \end{bmatrix} = \begin{bmatrix} k_1 & \dots & k_3 \end{bmatrix} \begin{bmatrix} w_{11} & w_{12} \\ \vdots & \vdots \\ w_{31} & w_{32} \end{bmatrix} + \begin{bmatrix} b_{01} & b_{02} \end{bmatrix}$$

DOT PRODUCT

MS

LINEAR CLASSIFIERS



If a dataset is
separable by a line
it's LINEAR SEPARABLE

LINEAR CLASSIFIER
↳ threshold functions

Sigmoid Function

$$f(z) = \frac{1}{1 + e^{-z}}$$

The diagram illustrates the components of logistic regression:

- Model:** $z = wX + b$
- Sigmoid:** $\sigma(z)$
- Prediction:** \hat{y}

The flow shows the model output z passing through the sigmoid function $\sigma(z)$ to produce the prediction \hat{y} .

LOGISTIC FUNCTION

```
import torch  
import torch.mn as mn  
import matplotlib.pyplot as plt
```

$$z = \text{torch.arange}(-100, 100, 0.1).view(-1, 1) \rightarrow z = \begin{bmatrix} -100 \\ \vdots \\ 100 \end{bmatrix}$$

$$\text{sig} = \text{m}. \text{Sigmoid}(\text{x}) \quad \Rightarrow \quad \sigma(x) = \begin{bmatrix} 0 & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & 1 \end{bmatrix}$$

$$\hat{y} = \text{sig}(z)$$

Final Project: READING COMPREHENSION

model = nn.Sequential
(nn.Linear(1, 1),
nn.Sigmoid())

$$\hat{y} = \text{model}(x)$$

CUSTOM MODE

```

class LogisticRegression(m.Module):
    def __init__(self, in_size, n):
        super(LogisticRegression, self).__init__()
        self.linear = m.Linear(in_size, n)
    def forward(self, x):
        return torch.sigmoid(self.linear(x))

```

MAKING PREDICTIONS

`torch.tensor([1 3 2 3])`

\hat{y} = $\text{model}(x)$ ← predict

-019-

BERNOULLI DISTRIBUTION

θ = Bernoulli parameter

$$\begin{array}{ccc} \text{COIN} & \text{FLIP} & \\ \text{HEAD} & \text{HEAD} & \text{TAILS} \\ \theta & \times & \theta \times (1-\theta) \\ 0.2 & \times & 0.2 \times 0.8 = 0.08 \end{array}$$

BERNOULLI DISTRIBUTION

$$\theta = \arg\max(p(Y|\theta))$$

each probability is a function

↳ $\ln(p(Y|\theta))$ EASY TO DEAL WITH

$$l(\theta) = \ln(p(Y|\theta)) = \sum_{m=1}^N y_m \ln(\theta) + (1-y_m) \ln(1-\theta)$$

	HEAD	TAIL	HEAD	TAIL	LIKELIHOOD
$\theta=0.5$	0.5	0.5	0.5	0.5	0.0625
$\theta=0.2$	0.2	0.8	0.2	0.8	0.0256

We can estimate the actual parameter by considering parameter values that maximize our likelihood

CROSS ENTROPY LOSS

loss function

$$l(\theta) = \frac{1}{N} \sum_{m=1}^N (y_m - \sigma(x_m+b))^2$$

we treated

$$l(\theta) = \sum_{m=1}^N (y_m - \text{THR}(x_m+b))^2$$

CROSS ENTROPY LOSS
WE SIGMOID NO HORIZONTAL PROBLEM

HSE

MAY NOT REACH MINIMUM

MAXIMUM LIKELIHOOD

↳ obtain a parameter that maximize function

CROSS ENTROPY (COST)

$$l(\theta) = -\frac{1}{N} \sum_{m=1}^N y_m \ln(\sigma(wx_m+b)) + (1-y_m) \ln(1-\sigma(wx_m+b))$$

↳ pytorch implementation

def criterion(yhat, y):

```
out = -1 * torch.mean(y * torch.log(yhat) + (1-y) * torch.log(1-yhat))
return out
```

MEAN SQUARE ERROR
criterion = nn.MSELoss()

BCE Loss()
for cross entropy

LOGISTIC REGRESSION in pytorch

for epoch in range(100)

for (x,y) in trainloader

yhat = model(x)

loss = Criterion(yhat,y)

optimizer.zero_grad()

loss.backward()

optimizer.step()

DEEP LEARNING PYTORCH

M1

SOFTMAX input classification

```
import torch.nn as nn
class Softmax(nn.Module):
    def __init__(self, in_size, out_size):
        super(Softmax, self).init()
        self.linear = nn.Linear(in_size, out_size)

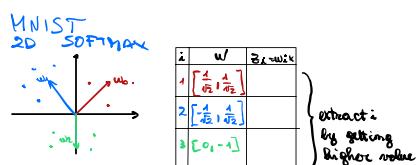
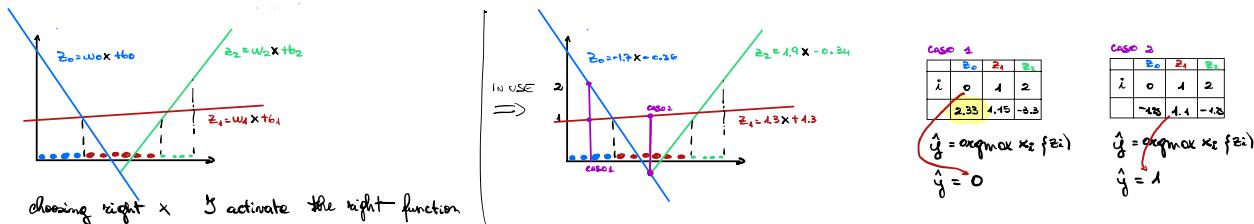
    def forward(self, x):
        out = self.linear(x)
        return out
```

from torch import Linear
torch.manual_seed(1)

```
model = Softmax(2,3)
x = torch.tensor([1.0, 2.0])
z = model(x)
yhat = z.max(1)
```

MULTIDIM INPUT
x = torch.tensor([1.0, 1.0], [1.0, 2.0], [1.0, 3.0])
z = model(x)
yhat = z.max(1)

SOFTMAX USING LINES TO CLASSIFY DATA



QUIZ

given an array of tensors
yhat = $z \cdot \text{max}(1)$
↳ extract for each array index of max value

Softmax(input features, # of classes)

CROSSENTROPY LOSS

SOFTMAX PyTorch

transforms.ToTensor()

↳ CONVERTS THE IMAGE TO A TENSOR

x.view(-1, 28*28)

↳ IMAGE ARE SQUARES SO WE HAVE TO CONVERT THEM TO VECTOR

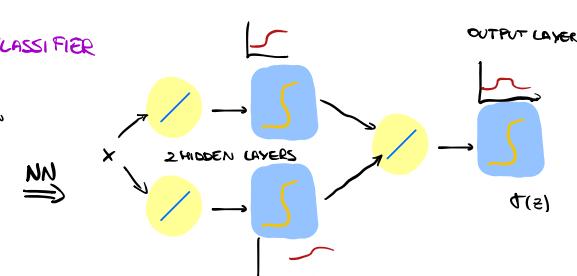
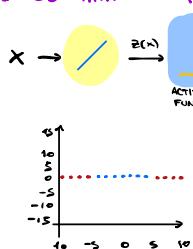
image is 10x10, converting to vector
how many elements does it have?

↳ 100

NEURAL NETWORK

is a function that can be used to approximate most functions

BUILD nn WITH LINEAR CLASSIFIER



PYTORCH IMPLEMENTATION

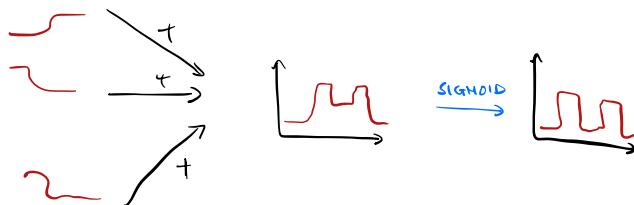
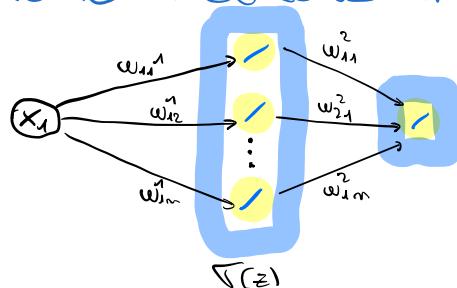
```
import torch
import torch.nn as nn
from torch import sigmoid
```

```
class Net(nn.Module):
    def __init__(self, D_in, H, D_out):
        super(Net, self).__init__()
        self.linear1 = nn.Linear(D_in, H)
        self.linear2 = nn.Linear(H, D_out)
    def forward(self, x):
        x = sigmoid(self.linear1(x))
        x = sigmoid(self.linear2(x))
        return x
```

```
model = Net(1, 2, 1)
x = torch.tensor([0.0])
yhat = model(x)
```

=>

MORE NEURONS IN HIDDEN LAYER

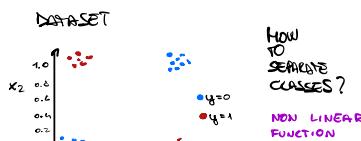
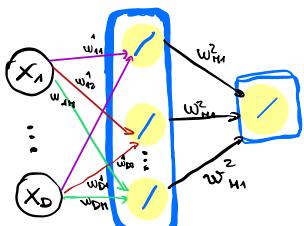


CODE

```
model = Net(1, 6, 1)
optimizer = torch.optim.Adam(model.parameters(), lr=0.01)
train(data_set, criterion, train_loader, epochs=1000)
```

```
→ model = torch.nn.Sequential(
    torch.nn.Linear(1, 7),
    torch.nn.Sigmoid(),
    torch.nn.Linear(7, 1),
    torch.nn.Sigmoid()
)
```

NEURAL NETWORKS WITH MULTIPLE DIMENSION INPUT



CODE

```
import torch
import torch.nn as nn
from torch import sigmoid
from torch.utils.data import Dataset, DataLoader
```

```
# Train Function
def train(data_set, model, criterion, train_loader, optimizer, epochs):
    cost = []
    acc = []
    for epoch in range(epochs):
        total = 0
        for x, y in train_loader:
            optimizer.zero_grad()
            yhat = model(x)
            loss = criterion(yhat, y)
            loss.backward()
            loss.backward()
            optimizer.step()
```

```
class XOR_Data(Dataset):
    # Constructor
    def __init__(self, N_s=100):
        self.x = torch.zeros(N_s, 2)
        self.y = torch.zeros(N_s, 1)
        for i in range(N_s // 4):
            self.x[i, :] = torch.Tensor([0.0, 0.0])
            self.y[i, 0] = torch.Tensor([0.0])

            self.x[i + N_s // 4, :] = torch.Tensor([0.0, 1.0])
            self.y[i + N_s // 4, 0] = torch.Tensor([1.0])

            self.x[i + N_s // 2, :] = torch.Tensor([1.0, 0.0])
            self.y[i + N_s // 2, 0] = torch.Tensor([1.0])

            self.x[i + 3 * N_s // 4, :] = torch.Tensor([1.0, 1.0])
            self.y[i + 3 * N_s // 4, 0] = torch.Tensor([0.0])

        self.x = self.x + 0.01 * torch.randn(N_s, 2)
        self.len = N_s

    # Getter
    def __getitem__(self, index):
        return self.x[index], self.y[index]

    # Get Length
    def __len__(self):
        return self.len
```

Model Creation

```
class Net(nn.Module):
    def __init__(self, D_in, H, D_out):
        super(Net, self).__init__()
        # Hidden Layer
        self.linear1 = nn.Linear(D_in, H)
        # Output Layer
        self.linear2 = nn.Sigmoid(H, D_out)

    def forward(self, x):
        x = torch.sigmoid(self.linear1(x))
        x = torch.sigmoid(self.linear2(x))
        return x
```

```

# cumulative loss
total += loss.item()
acc.append(accuracy(model,data_x))
cost.append(total)
return cost

```

TRAINING PROCESS

```

criterion = nn.BCELoss()
data_set = XOR_DATA()
train_loader = Data Loader(dataset = data_set, batch_size=1)

```

```

model = Net(2,4,1)
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
train(data_set, criterion, train_loader, optimizer, epochs=500)

```

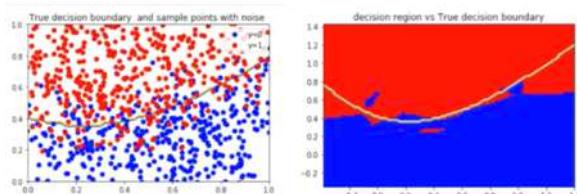
OVERFITTING

model too complex for the data
 → TOO MANY NEURONS IN THE HIDDEN LAYER

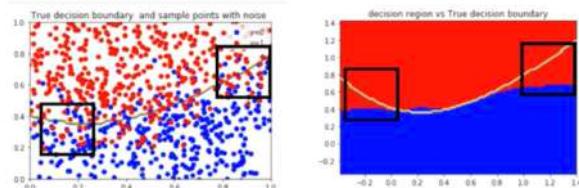
UNDERFITTING

model cannot capture the complexity of the data
 → TOO FEW NEURONS

Example of Overfitting



Example of Underfitting

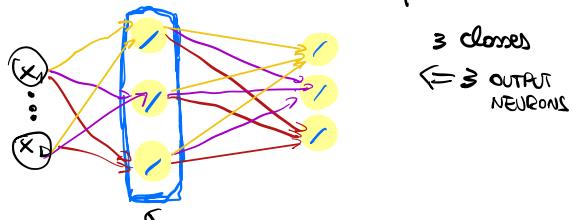


↳ SOLUTION

- USE VALIDATION DATA TO DETERMINE OPTIMUM # OF NEURONS
- GET MORE DATA
- REGULARIZATION: DROPOUT

MULTI CLASS NEURAL NETWORK

same number as # of classes



3 classes
 $\Leftrightarrow 3$ OUTPUT NEURONS

in PyTorch

```

class Net(nn.Module):
    # Constructor
    def __init__(self,D_in,H,D_out):
        super(Net,self).__init__()
        self.linear1=nn.Linear(D_in,H)
        self.linear2=nn.Linear(H,D_out)

    # Prediction
    def forward(self,x):
        x=torch.sigmoid(self.linear1(x))
        x=self.linear2(x)
        return x

```

With Sequential

```

model = torch.nn.Sequential(
    torch.nn.Linear(input_dim,hidden_dim),
    torch.nn.Sigmoid(),
    torch.nn.Linear(hidden_dim,output_dim))

```

Quiz

1. correct cost function for Multi Class NN? → `nn.CrossEntropyLoss()`
2. What cost function can be used for a 2 class problem? → `nn.BCELoss()`
3. The prediction step in a multi-class network utilizes the same procedure as softmax function? → `True`

BACK PROPAGATION

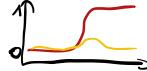
Reduces computation

USES THE DERIVATIVE OF THE 1st PARAMETER IN THE OUTPUT LAYER TO HELP CALCULATE THE PARAMETER OF THE NEXT LAYER.

! VANISHING GRADIENT → SOLUTION
change activation function
or optimizator methods 2.6

ACTIVATION FUNCTIONS

SIGMOID



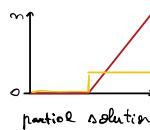
! VANISHING GRADIENT

TANH



! VANISHING GRADIENT

RELU



partial solution

```
class Net(nn.Module):
    def __init__(self,D_in,H,D_out):
        super(Net,self).__init__()
        self.linear1=nn.Linear(D_in,H)
        self.linear2=nn.Linear(H,D_out)
    def forward(self,x):
        x=torch.sigmoid(self.linear1(x))
        x=self.linear2(x)
        return x
```

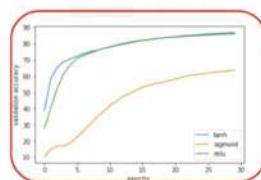
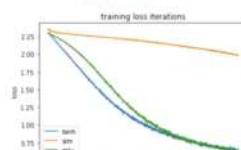
```
class Net_Tanh(nn.Module):
    def __init__(self,D_in,H,D_out):
        super(Net_Tanh,self).__init__()
        self.linear1=nn.Linear(D_in,H)
        self.linear2=nn.Linear(H,D_out)
    def forward(self,x):
        x=torch.tanh (self.linear1(x))
        x=self.linear2(x)
        return x
```

```
class NetRelu(nn.Module):
    def __init__(self,D_in,H,D_out):
        super(NetRelu,self).__init__()
        self.linear1=nn.Linear(D_in,H)
        self.linear2=nn.Linear(H,D_out)
    def forward(self,x):
        x=torch.relu(self.linear1(x))
        x=self.linear2(x)
        return x
```

```
model_Tanh=nn.Sequential( torch.nn.Linear(input_dim, hidden_dim),
nn.Tanh(),nn.Linear(hidden_dim, output_dim))
```

```
modelRelu=torch.nn.Sequential( torch.nn.Linear(input_dim,hidden_dim),
nn.ReLU(),nn.Linear(hidden_dim, output_dim))
```

Results MNIST



H3

DEEP NEURAL NETWORKS

multi hidden layers

```
class Net(nn.Module):
```

```
    def __init__(self,D_in,H1,H2,D_out):
        super(Net,self).__init__()
        self.linear1=nn.Linear(D_in,H1)
        self.linear2=nn.Linear(H1,H2)
        self.linear3=nn.Linear(H2,D_out)
    def forward(self,x):
        x=torch.sigmoid(self.linear1(x))
        x=torch.sigmoid(self.linear2(x))
        x=self.linear3(x)
        return x
```



$\text{model} = \text{Net} (\text{D_in} = D, \text{H1} = m, \text{H2} = h, \text{D_out} = f)$

AUTOMATE with nn. ModuleList()

↳ CODE

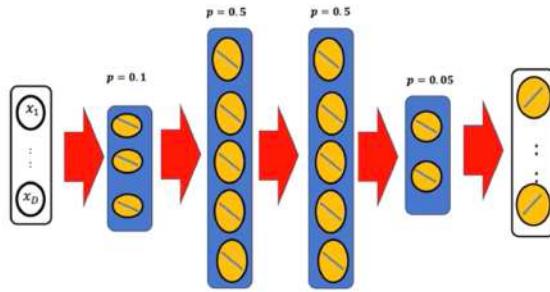
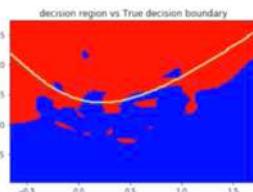
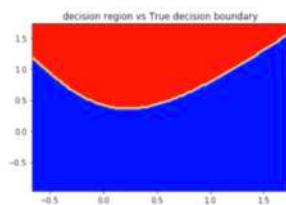
Layers = [2, 3, 4, 3]
 feature size
 neurons in layer
 to be neurons in next layer
 end layer

With Sequential

```
model = torch.nn.Sequential(
torch.nn.Linear(input_dim,hidden_dim1), torch.nn.Sigmoid(),
torch.nn.Linear(hidden_dim1, hidden_dim2), torch.nn.Sigmoid(),
torch.nn.Linear(hidden_dim2,output_dim1))
```

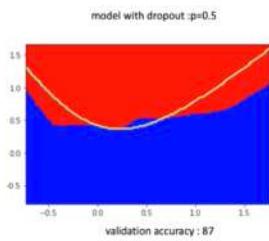
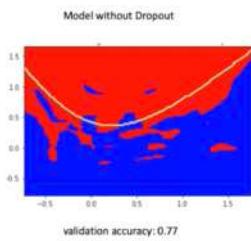
DROPOUT

$r \sim \text{Ber}(1-p)$



P is an Hyperparameter

EVALUATION



```
class Net(nn.Module):
    def __init__(self,in_size,n_hidden,out_size p=0):
        super(Net,self).__init__()
        self.drop=nn.Dropout(p=p) ←
        self.linear1=nn.Linear(in_size,n_hidden)
        self.linear2=nn.Linear(n_hidden,n_hidden)
        self.linear3=nn.Linear(n_hidden,out_size)
    def forward(self,x):
        x=torch.relu(self.linear1(x))
        x=self.drop(x) ←
        x=torch.relu(self.linear2(x))
        x=self.drop(x) ←
        x=self.linear3(x)
        return x
```

With Sequential

```
model= torch.nn.Sequential(torch.nn.Linear(1,
10),
torch.nn.Dropout(0.5),
torch.nn.ReLU(),
torch.nn.Linear(10,12),
torch.nn.Dropout(0.5),
torch.nn.ReLU(),
torch.nn.Linear(12, 1),
```

model_dropout = Net(2, 300, 2, p=0.5)
model_dropout.train()

optimizer = torch.optim.Adam(model_dropout.parameters(), lr=0.01)

! model.eval() NOT USE DROPOUT

QUIZ

- In what situation would you use dropout for classification? if training accuracy much larger test accuracy
- What is the purpose of using dropout? Reduce the impact of noise or overfitting
- Prediction w model w dropout \rightarrow model.eval() $yhat = model(x)$
- 60% of activations drop off \rightarrow nn.Dropout(0.4)

NEURAL NETWORK INITIALIZATION WEIGHTS

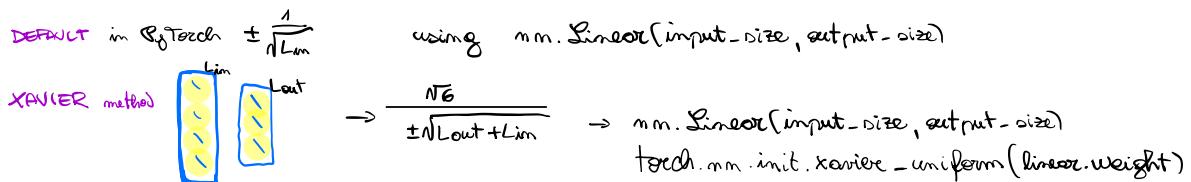
What happen when the weights in the same layer have the same values

WHEN WE DON'T INITIALIZE THE WEIGHTS PYTORCH RANDOMIZE FOR US

HOW TO FIX THE PROBLEM OF VANISHING GRADIENTS DUE TO A LARGE NUMBER OF NEURONS

\rightarrow Scale the width of the distribution by the inverse of the number of neurons. $\frac{1}{\# \text{of neurons}}$

DIFFERENT INITIALIZATION in PyTorch using Default Method



HE method $\rightarrow \text{linear} = nn.Linear(\text{input_size}, \text{output_size})$
`torch.nn.init.kaiming_uniform_(linear.weight, nonlinearity='relu')`

QUIZ

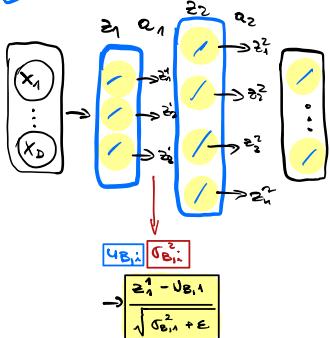
4. Why do we randomly initialize our parameters, instead of setting all the parameters to the same value?
 Each neuron will have the same output and such have the same gradient update

GRADIENT DESCENT with MOMENTUM

High velocity to go out of flat zone
 ↳ allows sometimes to go to the local minima faster

CODE: `torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)`

BATCH NORMALIZATION



It normalizes the distribution of inputs of each layer making the network less dependent on specific neuron weights. It operates on each batch of data. IMPROVE SPEED, PERFORMANCE & STABILITY OF NN

```
class NetBatchNorm(nn.Module):
    def __init__(self, n_size, n_hidden1, n_hidden2, out_size):
        super(NetBatchNorm, self).__init__()
        self.linear1=nn.Linear(n_size,n_hidden1)
        self.linear2=nn.Linear(n_hidden1,n_hidden2)
        self.linear3=nn.Linear(n_hidden2,out_size)

        self.bn1 = nn.BatchNorm1d(n_hidden1)
        self.bn2 = nn.BatchNorm1d(n_hidden2)

    def forward(self,x):
        x=F.sigmoid(self.bn1(self.linear1(x)))
        x=F.sigmoid(self.bn2(self.linear2(x))))
        x=self.linear3(x)
        return x
```

M4

CNN (convolutional NEURAL NETWORK)

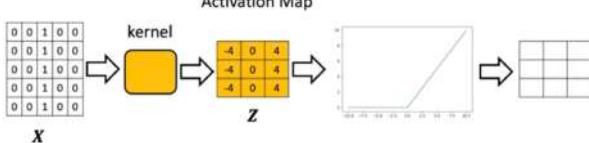
Looking at relative positions of the pixels rather than absolute positions.
 Using the image and something called the KERNEL, we are going to perform a convolution.
 The result is called an activation map.

CODE

```
conv=nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=2)
image=torch.zeros(1,1,5,5)
image[0,0,:]=1
z=conv(image)
```

Annotations: $(H-W)/\text{stride}+1$, $(H-W)/\text{kernel size} + 1$

ACTIVATION FUNCTIONS & MAX POOLING → will reduce the size of the activation maps reducing the number of parameters



MULTIPLE INPUT & OUTPUT CHANNELS

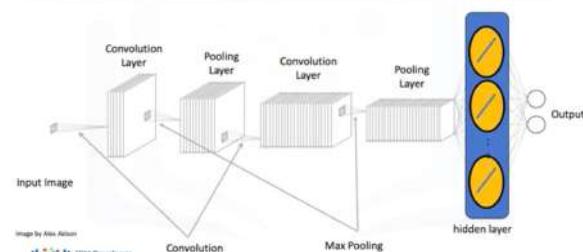
CASE: MULTIPLE OUTPUT

```
conv1 = nn.Conv2d(in_channels=1, out_channels=3, kernel_size=3)
image1 = torch.zeros(1, 1, 5, 5)
image2[0, 0, :, 2] = 1
z = conv1(image1)
```

Kernels = in * out - channels

CONVOLUTIONAL NEURAL NETWORK

CNN Architecture

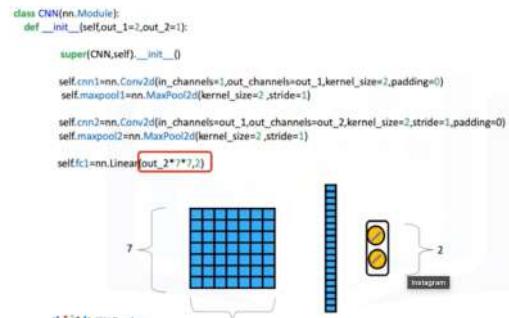


→ FORWARD METHOD

```
class CNN(nn.Module):
    def __init__(self, out_1=1, out_2=3):
        super(CNN, self).__init__()
        self.cnn1 = nn.Conv2d(in_channels=1, out_channels=out_1, kernel_size=5, padding=2)
        self.maxpool1 = nn.MaxPool2d(kernel_size=2)

        self.cnn2 = nn.Conv2d(in_channels=out_1, out_channels=out_2, kernel_size=5, stride=1, padding=2)
        self.maxpool2 = nn.MaxPool2d(kernel_size=2)
        self.fc1 = nn.Linear(out_2 * 4 * 4, 10)
```

The code defines a CNN class with two convolutional layers and a fully connected layer. The forward method processes the input through these layers, applying activation and pooling operations at each step. The final output is a 10-dimensional vector.



TORCH VISION MODELS

ResNet18 use pretrained models

Pretrained

```
model = models.resnet18(pretrained=True)
mean = [0.485, 0.456, 0.406]
std = [0.229, 0.224, 0.225]

composed= transforms.Compose([transforms.Resize(224),
    transforms.ToTensor(),
    transforms.Normalize(mean, std)])

train_dataset=Dataset(transform=composed, train=True )
validation_data=Dataset(transform=composed)
```

PyTorch

Creating the CNN

GPU

GPUs and Tensors

```
torch.tensor([1, 2, 3, 4])
torch.tensor([1, 2, 3, 4]).to(device)
tensor([1, 2, 3, 4], device='cuda:0')
```

Training

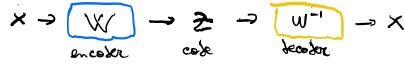
```
for epoch in range(num_epochs):
    for features, labels in train_loader:
        features, labels = features.to(device), labels.to(device)
        optimizer.zero_grad()
        predictions = model(features)
        loss = criterion(predictions, labels)
        loss.backward()
        optimizer.step()
```

```
model = CNN()
model.to(device)
```

45

AUTOENCODERS & MATRICES

MATRIX TRANSFORMATION



```
class AutoEncoder(nn.Module):
    def __init__(self, input_dim=2, encoding_dim=2):
        super(AutoEncoder, self).__init__()
        self.encoder = nn.Linear(input_dim, encoding_dim, bias=False)
        self.decoder = nn.Linear(encoding_dim, input_dim, bias=False)

    def forward(self, x):
        x=self.encoder(x)
        x=self.decoder(x)
        return x
```



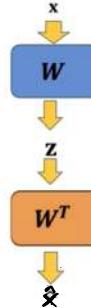
encoder & decoder are Matrix

```
auto_encoder_2Dcode=AutoEncoder(2,1)

W=torch.tensor([[1./2**0.5,1./2**0.5]])
auto_encoder_2Dcode.state_dict()['encoder.weight'].data[:,]=W
auto_encoder_2Dcode.state_dict()['decoder.weight'].data[:,]=torch.transpose(W,0,1)

z=auto_encoder_1Dcode.encoder(torch.tensor([1,0,1,0]))
z:tensor([0.3563]), grad_fn=<MmBackward>

x_hat=auto_encoder_1Dcode.decoder(z)
x_hat:tensor([-0.1929, 0.1561]), grad_fn=<SqueezeBackward3>
```



PCA Principal components analysis

$$X = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$$

$$x = x - x.\text{mean(dim=0)}$$

$$X = (x - \mu) D^{-1}$$

code: (before training network)

standard = preprocessing. StandardScaler

x_train = standard.fit_transform(x_train)

x_test = standard.transform(x_test)

PCA TO AUTOENCODERS

input D code E output

SHALLOW AUTOENCODERS

DEEP AUTOENCODERS

more than 1 layer w/ activation function
→ ANOMALY DETECTION

TENSORFLOW

M1 Data flow graph

NODES = Mathematical Operations
EDGES = Multi-Dimensional Array (tensors)

build \Rightarrow execute

Tensor multidimensional array

placeholders created before usage

TF 2.X & EAGER EXECUTION

Keras is high level API

Performance OPT \rightarrow GPU

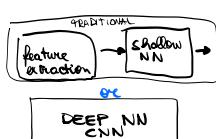
<pre>import tensorflow as tf import numpy as np a = tf.constant(np.array([1., 2., 3.])) b = tf.constant(np.array([4., 5., 6.])) c = tf.matmul(a, b) session = tf.Session() output = session.run(c) session.close()</pre> <p>OLD</p>	<p>NEW</p> <pre>import tensorflow as tf import numpy as np a = tf.constant(np.array([1., 2., 3.])) b = tf.constant(np.array([4., 5., 6.])) c = tf.matmul(a, b) output = c.numpy()</pre> <p>eager execution</p>
--	--

DEEP LEARNING

DEEP NEURAL NETWORKS

CNN \rightarrow Image & Vision

OBJECT DETECTION



RNN \rightarrow Recurrent Neural Network

- \rightarrow sentiment
- \rightarrow next word suggestion
- \rightarrow translation
- \rightarrow speech to text

RBH Restricted Boltzmann Machines

- \rightarrow RECONSTRUCT DATA
- \rightarrow feature extraction / learning
- \rightarrow dimensionality reduction
- \rightarrow pattern recognition
- \rightarrow recommender systems
- \rightarrow handling missing values
- \rightarrow topic modelling

AUTOENCODER

- \rightarrow UNSUPERVISED
- dimensionality reduction
- feature extraction
- image recognition

DBN Deep Belief Network

- \rightarrow BACK PROP. PROBLEM
- \rightarrow CLASSIFICATION very ACCURATE
SMALL DATASET

M2

CNN

Key features

- Detect & classify
- Independence from pose, scale, illumination, conformation & clutter

SEQUENCE OF STEPS

input image \rightarrow find primitive features \rightarrow detect object path \rightarrow predict element

CNN FOR CLASSIFICATION

feature extraction, feature selection

IMG \rightarrow AUTOMATIC
feature extractor \rightarrow
feature selector

SHALLOW NN \rightarrow multiple CNN

CONVOLUTIONAL POOLING LAYER FULLY CONNECTED LAYER

CNN ARCHITECTURE

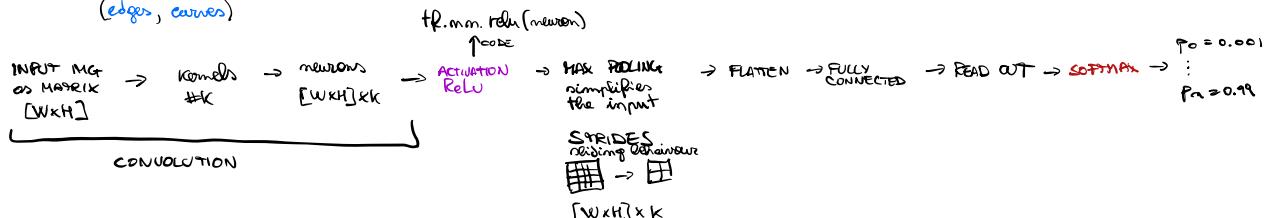
C1 CONVOLUTIONAL LAYER

↳ detect pattern & features from img EX: EDGES

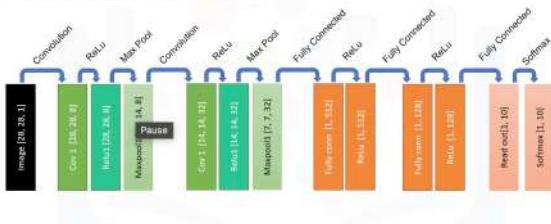
MATH → convolution w/ kernel

USING EXISTING FILTERS

(edges, curves)



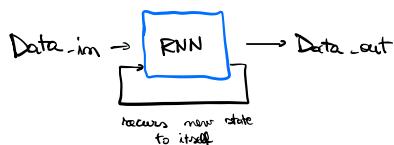
CNN architecture



M3

THE SEQUENTIAL PROBLEM

→ not handled well by traditional NN
• We want to have some MEMORY

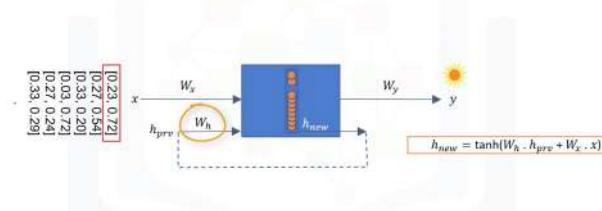


ISSUES

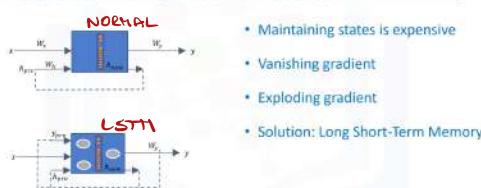
- Must remember all states at any given time
 - computationally expensive
 - only store states within a time window
- Sensitive to changes in their parameters
- Vanishing gradient (to 0)
- Exploding gradient (exp to ∞)

LSTM (Long Short Term Model)

The Recurrent Model

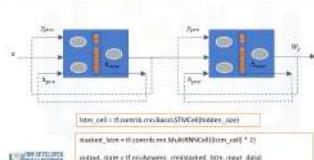


Recurrent Network Problems

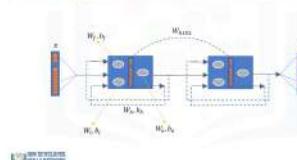


- Maintaining states is expensive
- Vanishing gradient
- Exploding gradient
- Solution: Long Short-Term Memory

Stacked LSTM



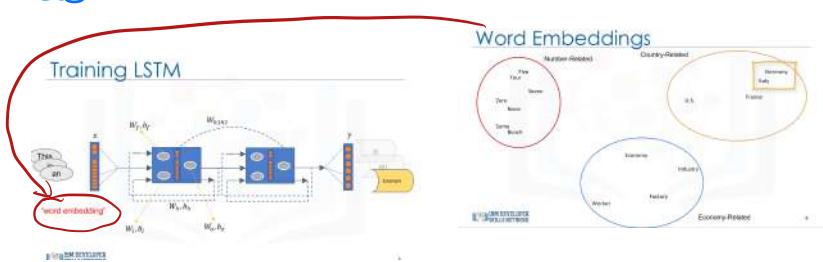
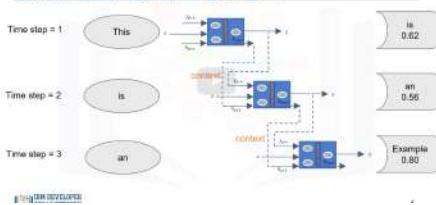
Training LSTM



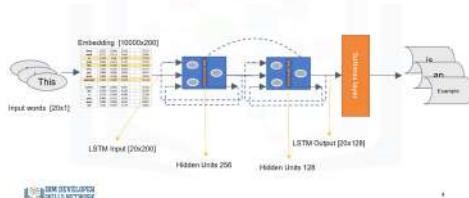
↳ code: `lstm_cell = tf.contrib.rnn.BasicLSTMCell(hidden_size)`

APPLYING RNN TO LANGUAGE MODELS

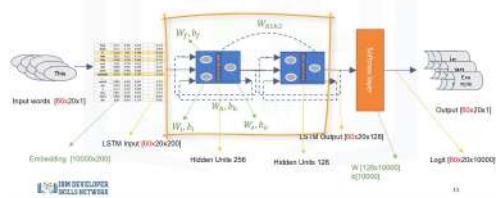
Unfolded LSTM network



Training LSTM



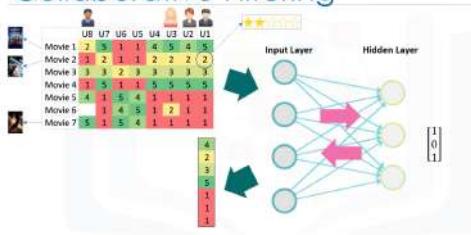
Training LSTM



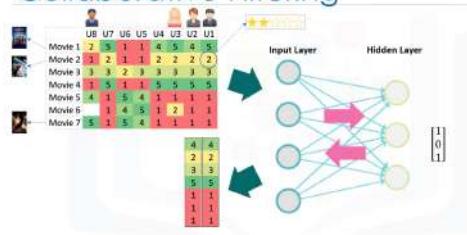
M4

RBM

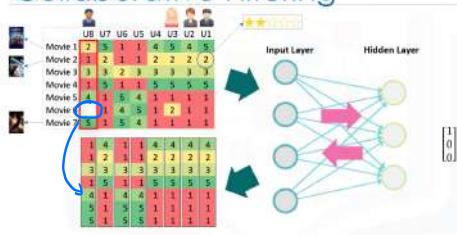
Collaborative Filtering



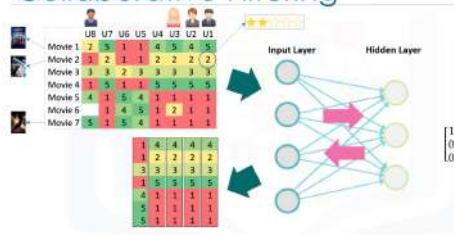
Collaborative Filtering



Collaborative Filtering



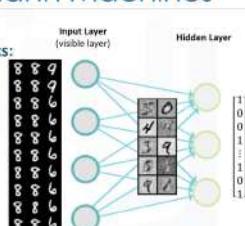
Collaborative Filtering



Restricted Boltzmann Machines

RBMs are shallow neural networks:

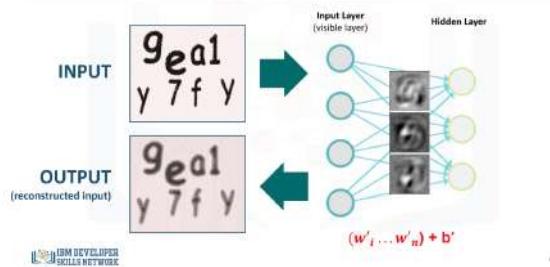
- 2 layers
- Unsupervised
- Find patterns in data by reconstructing the input



USAGE :

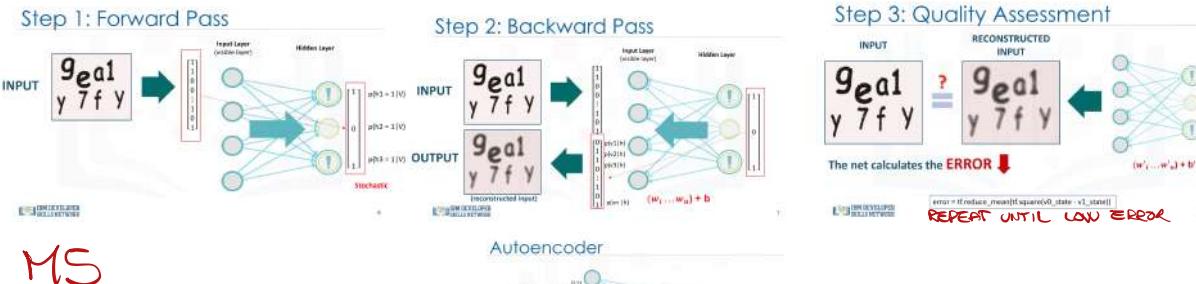
- dimensionality reduction
- Feature extraction
- Collaborative filtering
- main block of DBN

Learning Process of RBMs



ADVANTAGES of RBM

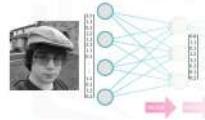
- good at handling UNLABELED data
- RBM extract important features from input
- RBM are more efficient at dimensionality reduction than PCA.



MS

AUTOENCODERS

- detect KEY FEATURES
- data compression
- Learning generative models of data
- Dimensionality reduction



- HIGH DIMENSION #
- time to fit in memory, exp
- SPARSITY
- OVERLAP
- also are
- PCA dimensionality reduction

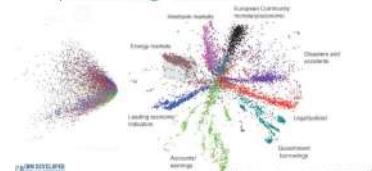
Curse of Dimensionality

$$m^{-p/(2p+d)}$$

Being:
m: Number of data points
d: Dimensionality of the data
p: Parameter that depends on the model

PCA ↔ AUTOENCODER
AUTOENCODER GIVES SEPARABILITY
→ clustering dogs available

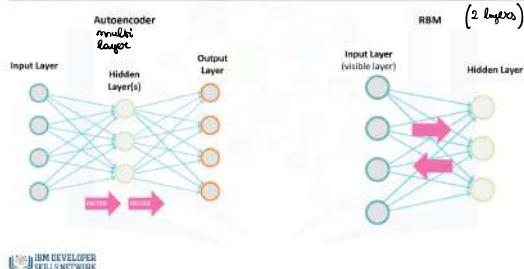
Comparison against PCA



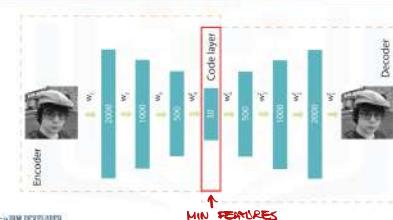
AUTOENCODER STRUCTURE

The main goal is to take unlabeled inputs encode them, and then try to reconstruct them afterwards based on the most valuable features identified in the data.

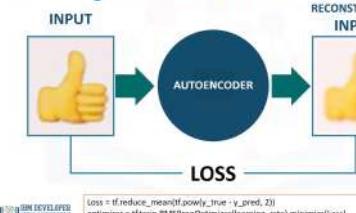
Autoencoder vs RBM



Autoencoder Architecture



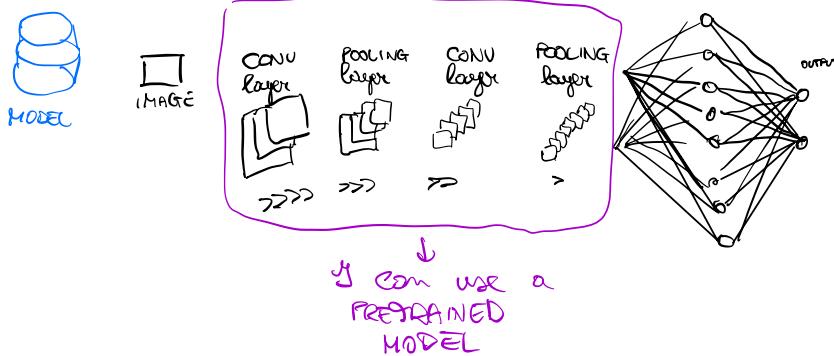
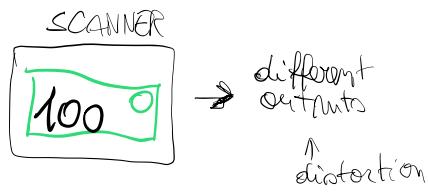
Learning Process of Autoencoders



CAPSTONE

- ATM

denomination	y
5	0
10	1
20	2
50	3
100	4
200	5
500	6



1. Loading Data

```
resources / data / training_data - pytorch
          / validation_data - pytorch
          / test_data - pytorch
```

```
import matplotlib.pyplot as plt
from PIL import Image

train_dir = 'resources / data / training_data - pytorch'
name = "0.jpg"

input = train_dir + name
img = Image.open(input)
plt.imshow(img)
plt.show()
```

2. Image Preprocessing

Given CSV file

#	TYPE	PATH	CLASS
0	S	0.jpg	0
1	S	1.jpg	0
2	10	2.jpg	1
:	10	3.jpg	1
n	:	n.jpg	5

class Dataset(Dataset):

def __init__(self, csv_file, data_dir, transform=None)

Load

- img file name

- class of each sample

- directories of image

def __get_item__(self, idx):

Load image files

Output img file

Load Class Name

```

INIT
    csv_file = "name_of_csv.csv"
    data_name = pd.read_csv(csv_file)
    self.data_name

    data_dir = "/keras/data/train.../test"
    self.data_dir

GET_ITEM
    file_name: train_data_name.iloc[0,2]
    y: data_name.iloc[0,3]

    img_name = data_dir + data_name.iloc[0,2]
    image = Image.open(img_name)
    y = data_name.iloc[0,3]
    if self.transform:
        image = self.transform(image)
    return image, y

```

3 TRAINING WITH PyTorch

We use ResNet18

last layer & OUTPUT layer

↳ CODE

```

model = models.resnet18(pretrained=True)
mean = [#, ... #]
std = [#, ..., #]

transform_stuff = transforms.Compose(
    transforms.Resize(224),
    transforms.ToTensor(),
    transforms.Normalize(mean, std))

train_dataset = Dataset(transform=composed, csv_file, data_dir=train...)
validation_dataset = Dataset(transform=composed, csv_file, data_dir=validat...)

for param in model.parameters():
    param.requires_grad = False

model.fc = nn.Linear(512, 7)

```

train_loader = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=15)

validation_loader = torch.utils.data.DATALOADER(dataset=validation_dataset, batch_size=15)

shuffle = True

criterion = nn.CrossEntropyLoss()

optimizer = torch.optim.Adam([param.. for param in model.parameters() if param.bk == 0.002])

N-EPOCHS = 100

loss_list = []

accuracy_list = []

correct = 0

n_test = len(validation_dataset)

```

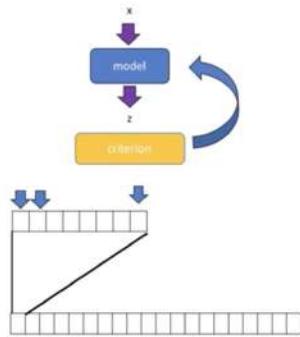
for epoch in range(n_epochs):
    loss_sublist = []
    for x, y in train_loader:
        model.train()
        optimizer.zero_grad()
        z = model(x)
        loss = criterion(z, y)
        loss_sublist.append(loss.data.item())
        loss.backward()
        optimizer.step()
    loss_list.append(np.mean(loss_sublist))

```

```

for x_test, y_test in validation_loader:
    z = model(x_test)
    yhat = torch.max(z.data, 1)
    correct += (yhat == y_test).sum().item()
accuracy = correct / n_test
accuracy_list.append(accuracy)

```



```
correct = 0
```

```
for x_test, y_test in validation_loader:
```

```
    model.eval()
```

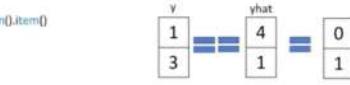
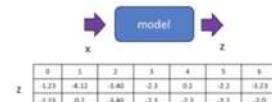
```
    z = model(x_test)
```

```
    ... yhat = torch.max(z.data, 1)
```

```
    correct += (yhat == y_test).sum().item()
```

```
accuracy = correct / n_test
```

```
accuracy_list.append(accuracy)
```



Manuel Oklizio

13.7.2023